CPU gross energy differences | MMLU−Pro

Five−repeat question summaries | A minus B | Unadjusted pointwise 95% intervals

Comparison: A − B	Estimate [95% interval]	Holm p	Holm decision
RQ1 Cascade − Large	−1,254.09 [−1,634.75, −873.42]	3.87e−09	Reject
RQ1 Cascade − Medium	1,530.39 [1,107.19, 1,953.59]	1.80e−10	Reject
RQ1 Cascade − Small	2,000.62 [1,562.69, 2,438.54]	2.58e−15	Reject
RQ1 MF − Large	−285.01 [−414.31, −155.71]	6.61e−05	Reject
RQ1 MF − Medium	2,499.47 [2,331.51, 2,667.43]	6.78e−73	Reject
RQ1 MF − Small	2,969.69 [2,784.93, 3,154.46]	2.40e−78	Reject
RQ2 MF − Cascade	969.08 [575.08, 1,363.08]	7.45e−06	Reject

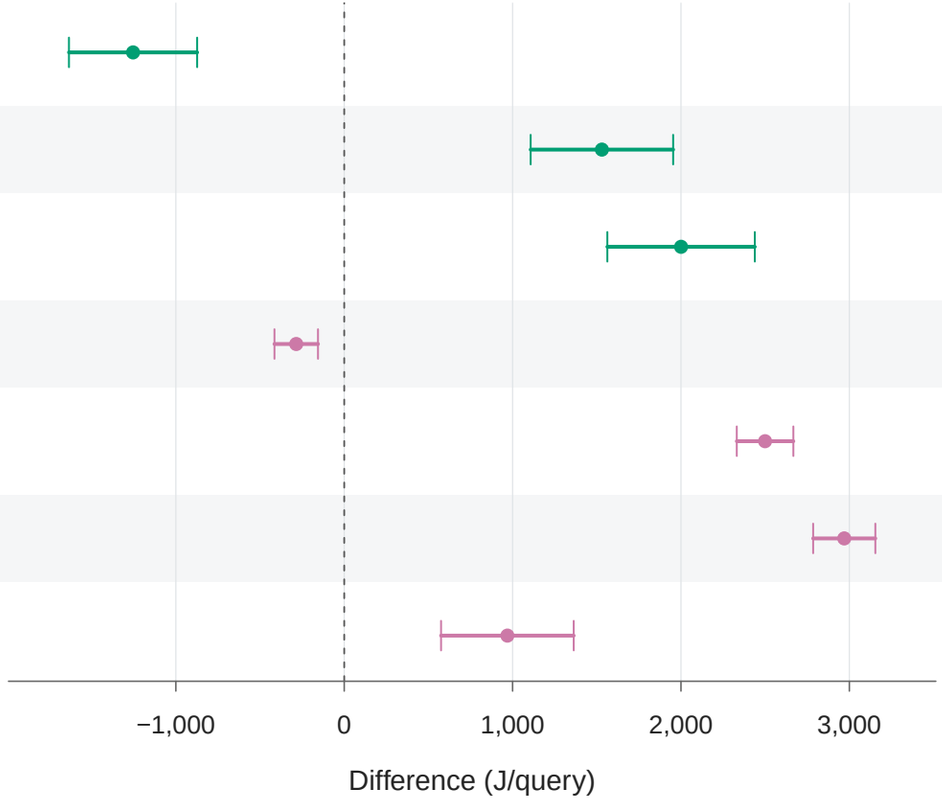

Intervals are unadjusted pointwise 95% intervals, not simultaneous intervals. Holm decisions use the original RQ families at alpha = 0.05.
Paired−t inference is conditional on the recorded runs and question model; it does not estimate uncertainty over future reruns.