GPU gross energy differences | MMLU−Pro

Five−repeat question summaries | A minus B | Unadjusted pointwise 95% intervals

Comparison: A − B	Estimate [95% interval]	Holm p	Holm decision
RQ1 Cascade − Large	−2,508.58 [−3,133.13, −1,884.03]	2.29e−12	Reject
RQ1 Cascade − Medium	2,680.40 [1,972.66, 3,388.15]	3.00e−11	Reject
RQ1 Cascade − Small	3,489.01 [2,755.57, 4,222.44]	2.43e−16	Reject
RQ1 MF − Large	−526.48 [−751.39, −301.56]	2.80e−05	Reject
RQ1 MF − Medium	4,662.51 [4,363.92, 4,961.10]	2.74e−76	Reject
RQ1 MF − Small	5,471.11 [5,144.21, 5,798.02]	2.89e−81	Reject
RQ2 MF − Cascade	1,982.10 [1,330.77, 2,633.44]	5.48e−08	Reject

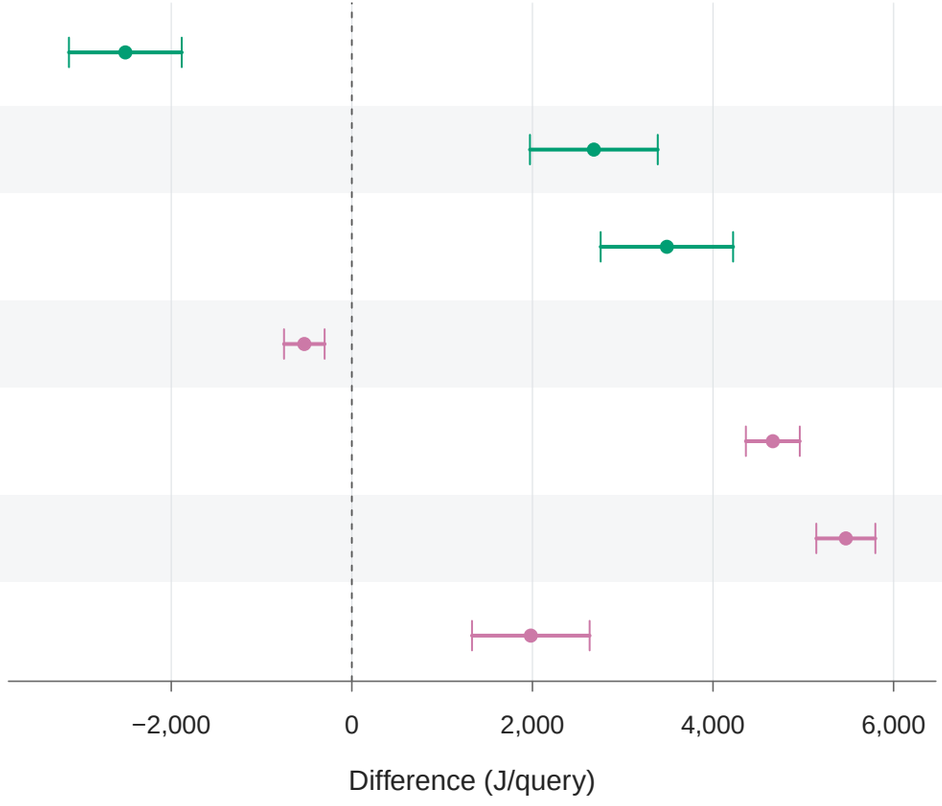

Intervals are unadjusted pointwise 95% intervals, not simultaneous intervals. Holm decisions use the original RQ families at alpha = 0.05.
Paired−t inference is conditional on the recorded runs and question model; it does not estimate uncertainty over future reruns.